Mongo DB – Week 2

Creating Documents:

→ Insert one ()

→ Insert Many () : Has option to do bulk inserts as well as insert ordering to skip failed records.

_id field → mongo creates one id field by default (Primary index) if User has not specified one.

ObjectId : DATE | MAC ADDR | PID | COUNTER
12-BYTE HEX STRING

Reading Documents:

* db.movieDetails.find ({rated : "PG-13"}).count()
(Query Document)
(Or) ⇒ 152 records

db.movieDetails.find ({rated : "PG-13", Year : 2009}).
⇒ 8 records. count()

⇒ db.movieDetails.find ({"tomato.meter" : 100}).pretty()
* → Nested Document find.

Equality Matches:

(1) On Entire Array:

Ex: db.movieDetails.find({"writers": ["Ethan Coen", "Joel Coen"]}).count()

* Here order of element matters.

(2) Based on any Element:

db.movieDetails.find({"actors": "Jeff Bridges"}).pretty()

* will match actors having Jeff Bridges in the movie list.

(3) Based on Specific Element:

Ex: find movies where Jeff Bridges is the main actor

db.movieDetails.find({"actors.0": "Jeff Bridges"}).pretty()

* > .0 means first element in actors array.

* find() method returns a cursor

So: var c = db.movieDetails.find

var doc = function() { return c.hasNext() ? c.objLeftInBatch() : c.next() : null; }

doc() → next Element using above function.

Projection: Reduce the size of data returned by a Query. (limit the fields returned)

Ex: db.movieDetails.find({rated:"PG"}, {title:3}).Pretty()

Will contain only

* Explicitly exclude some fields.

{title:1, _id:0}

This will exclude the _id field in results.

Query Selectors:

$eq, $gt, $gte, $Lt, $Lte

Ex: Find all movies having runtime greater than 90.

db.movieDetails.find({runtime: {$gt:90}}).Pretty()

* db.movieDetails.find({runtime: {$gte:90, $Lte:120}}).Pretty()

db.movieDetails.find({rated: {$ne:"UNRATED"}}).Count()

db.movieDetails.find({rated: {$in:["G","PG"]}}).Pretty()

$exists, $type → If a field is of specified type

Ex:

db.movieDetails.find({"tomato.meter": {$exists: true}}).count()

db.movieDetails.find({"_id": {$type: "String"}}).count()

Logical Operators:

$or, $and, $not, $nor (logical NOR returns all documents that fail to match both clauses)

db.movieDetails.find({$or: [{"tomato.meter": {$gt: 95}}, {"metacritic": {$gt: 88}}]}).pretty()

(or)

({$and: [{"tomato.meter": {$gt: 95}}, ...

* $and mostly used to specify multiple criteria on same field.

Regex Operators:

db.movieDetails.find({"awards.text": {$regex: /^won 1.*/}}

Array Operators:

$all → Matches arrays that contain all elements specified in the query.

Ex:

```
db.movieDetails.find({genres: {$all: ["Comedy", "Crime", "Drama"]}}).pretty()
```

```
db.movieDetails.find({countries: {$size: 1}}).pretty()
```

$elemMatch

```
db.movieDetails.find({boxOffice: {country: "UK", revenue: {$gt: 15}}})
```

↓

This would match the boxOffice element and look for country UK and revenue greater than 15 in different elements. So use below!

```
db.movieDetails.find({boxOffice: {$elemMatch: {country: "UK", revenue: {$gt: 15}}}})
```

UPDATING DOCUMENTS.

1. Update One
2. Update many
3. Replace one

```
db.movieDetails.updateOne({title: "The Martian"},
                          {$set: {poster: "http://....."
```

Update field Operators: $rename, $set, $unset

```
db.movieDetails.updateOne({title: "The Martian"},
                          {$inc: {"tomato.reviews": 3, "tomato.userReviews":
                                                              2533})
```

* For array field updates use $push.

$push along with $each.
$slice → keep only these many elements in the array. For example 5 more recent.
$position → Where to push this array element

* Update many → Update all documents that match the filter.

```
db.movieDetails.updateMany({rated: null},
                           {$unset: {rated: ""}})
```

UPSERTS : → like DB2 Merge.

Mongo DB Schema Design.

⇒ Application Driven Schema : 3rd Normal Form.

* Rich Documents.
* Pre Join / Embed Data
* No Joins in Mongo DB (Embed directly)
* No Constraints.
* Atomic Operations
* No Transaction Support.
* No Declared Schema.

* Avoid Constraints & go for embedding Data using Collections.

* Atomic Operations (i.e all or nothing) in mongo DB instead of Transactions.

One to one Relations

* Example : Employee ⟷ Resume

Come to have it in Separate Collection. (i.e to embed or not)
- → Frequency of Access
- → Size of items (Mongo DB doc 16 MB)
- → Atomicity of data

One to Many Relations.

Example : City to Persons.

(NY) → (8 million People)

blogPosts : Comments

(1) : (10)

(One Scenario) Posts

```
{ name
  Comments : [
  ]
}
```

vs People

```
{ name : "Syni"
  city : "NYC"
}
```

and

City.

```
{ _id : "NYC"
}
```

⇒ one to few Scenario.

Many to Many Relations :

Ex: Books : Authors
Students : teachers

⇒ Few : Few.

Authors

_id : 27
Author_name : "mitchel Stark"
books : [12, 7, 8]

Books.

_id : 12
title : "Gone with the Coind"

Multi key Indexes :

Students

{ _id : 0
name : 'Andrew',
teachers : [1, 7, 10, 23] }

teachers.

{ _id : 10,
name : 'Tony Stark' }

db.Students.find({'teachers' : {$all : [0,1]}})

And

db.Students.find({'teachers' : {$all : [0,1]}}).explain()

* Cursor : Btree Cursor
isMultikey : true.

* Const it did while the query got created.

Benefits of Embedding:

→ Improved Read Performance

→ One Round trip to DB

* 3NF will avoid Modification Anamolies.

ODM – Object Document Mappers.

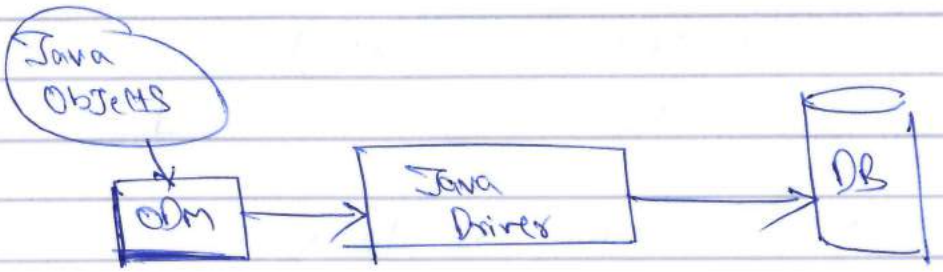

* Morphia → Java ODM for MongoDB.

MongoDB (week 4)

→ Pluggable Storage Engines.

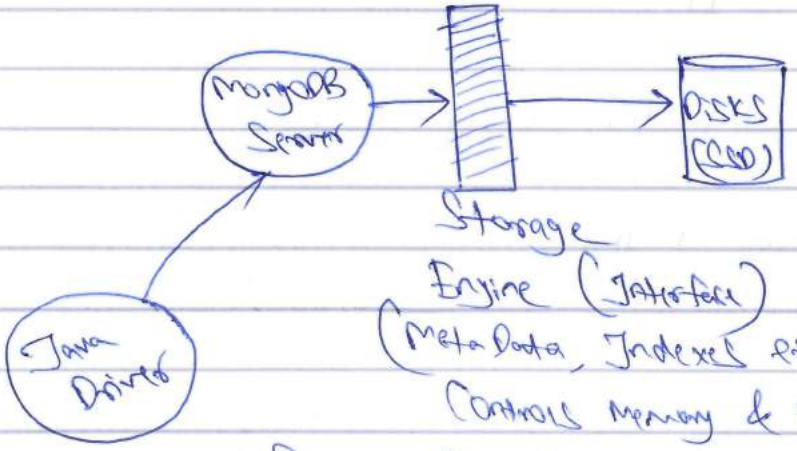

Storage Engines:

* MMAP → Collection level locking.
* Wired Tiger → Not Default.
  * Offers Document level Concurrency.
  * Compression of documents & indexes
  * No Inplace Update

Index

* Ordered list of things having Pointer to original Record on the Disk.
* BTree implementation.

Writes : Slower
Reads : Faster
} Indexes Disadvantage.

db.Students.CreateIndex({Student_id: 1}); → Ascending
-1 = Descending.

* db.Students.Explain().find({Student_id: 53});
Query Plan on the Collection followed by the Command. And winning Plan can be checked.

* Explain(true) ⇒ To check how many documents were Examined.

db.Students.getIndexes(); ⇒ To check the index on a Collection
↳ By Default index will be on _id
db.Students.dropIndex({Student_id: 1});

Multi key Indexes

Eg: { name : 'Andrew',
tags : ['Photography', 'hiking', 'golf'],
Color : 'Red',
location : ['NY', 'CA']}

Index on tags or (tags, Color)

But (tags, location) can't be index because both tags & location is an array document.

- Notation & multi key.

* db.students.createIndex({'scores.score':1});

Unique Indexes:

db.stuff.createIndex({thing:1}, {"unique":true});

* Sparse Indexes → when index key is missing from some of the documents.

For example:

{a:1, b:2, c:5}	Index on
{a:10, b:5, c:10}	a → possible
{a:13, b:17}	b → possible
{a:7, b:23}	c → not possible since c value not in some docs.

* This is when sparse index helps which won't include document 3 and 4 in index since it does not have c value.

⇒ db.employees.createIndex({cell:1}, {unique:true, sparse:true});

Sparse Index → uses collection scan, so don't use for sorting.

Background Indexes:

db.students.createIndex({'scores.score':1}, {background:true})

⇒ Slow
⇒ Don't block readers/writers.

MongoDB explain:

db.foo.explain().help()

* var exp = db.example.explain()
exp.help();
exp.find({a:17, b:553}).sort({b:-1});

Refined

Explain:

→ Query Planner (Default)
→ Execution Stats
→ allPlans Execution

```
var exp = db.example.explain("executionStats");
exp.find({a:17, b:55});
```

* Winning Plan + Rejected Plans + Detailed info regarding the several stages of Query execution.

```
var exp = db.example.explain("allPlansExecution");
```

* Execution Stats for all Plans.
⇒ allPlansExecution array of all plans.

Covered Indexes:

* exp.find({i:45, j:23}, {_id:0, i:1, j:1, k:1});
↳ Don't include _id in the display.

* exp.find({i:45, j:23}, {_id:0});
↳ more documents examined as mongoDB don't know how many more variables are there like i, j or k etc.

Choosing an Index:

Example Indexes: ① b, c
② c, b
d, e
e, f
③ a, b, c

⇒ Potential candidates so three winning plan and check which is fastest.

⇒ Winning Criteria
① Returned ok results
② Threshold values.
③ Cache winning query.

Threshold writes | Rebuild | Mongodb restart

db.collection.stats()
↳ will give totalIndexSize()
(or) db.collection.totalIndexSize()

★ So make sure your Indexes fits in the memory.

Index Cardinality.

Regular	Sparse	Multi key
1:1	≤ documents	tags: [—, —, —] > documents.

Geo Spatial Indexes.

⇒ Find things based on location.

i.e Search based on location.

* db.Collection.ensureIndex({"location": '2d', type: 1});
⇒ 2 Dimensional Index.
→ Ascending

db.Collection.find({location: {$near: [50, 50]}})
(or)
db.Collection.find({location: {$near: [50, 50]}}).limit(3);

GeoSpacial Spherical.

* 2d Sphere
* GeoJson for latitude, longitude Coordinates.

db.Collection.ensureIndex({'location': '2dsphere'})
⇓
for GeoJSON Document Index.

Example Query:

```
db.Places.find({
    location: {
        $near: {
            $geometry: {
                type: "Point",
                Coordinates: [-122.16, 37.42]},
                $maxDistance: 2000}
}) }
```

Full Text Search.

db.collection.ensureIndex({'words': 'text'}) → text Index.

Example: db.sentences.find({$text: {$search: 'dog'}})

textScore ⇒ Document with more word successfully matched.

Ex:

db.sentences.find({$text: {$search: 'dog tree Obsidian'}}, {score: {$meta: 'textScore'}}).sort({score: {$meta: 'textScore'}})

Designing Using Indexes.

Goal: Efficient Read/write Operations.

logging & Profiling:

⇒ logs slow queries by default for query running more than 100ms.

Profiler: Levels

0	1	2
Off	log Slow Queries	All my Queries.

* db.system.Profile.find({millis: {$gt:133}}).sort({ts:1}).
↳ Query which is greater than 1ms pretty(),
and sorted by time stamp ~~descending~~.
Ascending.

→ db.getProfilingLevel() ⇒ To get level of logging.
→ db.setProfiling~~Status~~Level(1,4)

Mongotop.

* MongoDB Profiler. With read & write time information.

Mongostat → What is going on during insert, update, delete queries.

SHARDING:

⇒ Splitting up large collection among multiple servers.

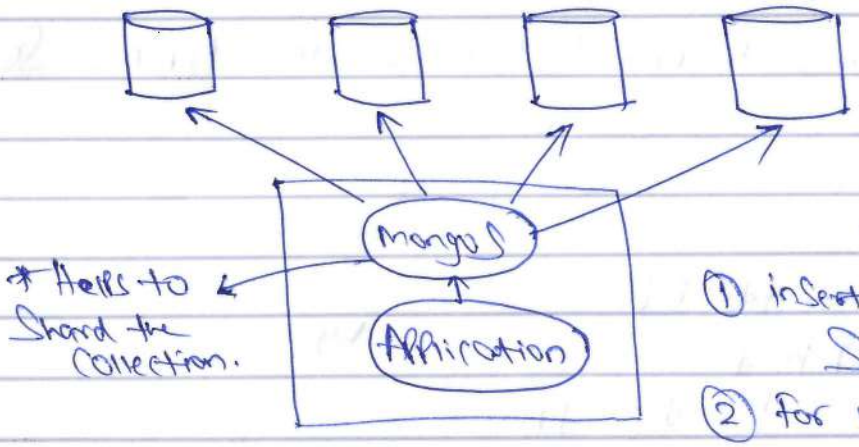

* Tells to Shard the Collection.

* Based on the Query send request to the right Server instance.

(1) insert Shard have a Shard key.

(2) for update / remove / find if Shard key is not there then it sends requests to all Servers (i.e broad casting)

Week 5 — Aggregation.

⇒ Similar to Count, GroupBy aggregation in RDBMS SQL.

MongoDB example:

```
db.Products.aggregate([
    { $group:
        { _id: "$manufacturer",
          num_products: { $sum: 1 }
        }
    }
])
```

(key)

Aggregation Pipeline:

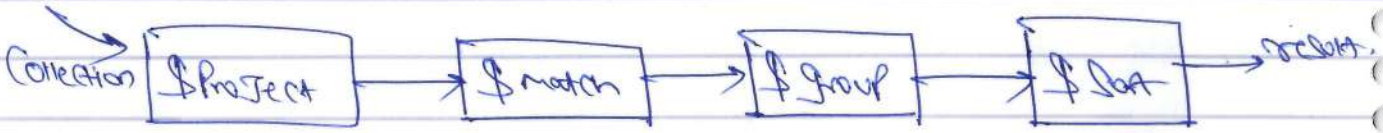

$Project — reshape a document — 1:1 (one document will have)
$match — filter stage — n:1
$group — aggregate — n:1 ($sum or $count)
$Sort — Sort — 1:1
$Skip — Skips — n:1
$limit — limits — n:1

$unwind — Flatten the data

Ex: Initial document tags: ["red", "blue", "green"]

⇓

tags: red, tags: blue, tags: green

Compound Grouping:

```
_id: { "manufacturer": "$manufacturer"
       "Category": "$category" },
num_products: { $sum: 1 }
```

★ _id field can be a document also. and not always a scalar value.

Aggregation Expressions Overview.

$sum, $avg, $min, $max, $push, $addToSet — Push to Array

$first, $last — Need to first sort the document Otherwise it will be Arbitrary.

$addToSet → Add to collection if its not there.

```
db.Products.aggregate([
    {$group:
        {_id: {"maker": "$manufacturer"},
         categories: {$addToSet: "$category"}
        }
    }])
```

maker wise find out all the categories.

Double Grouping:

```
db.fun.find()
    {"_id":0, "a":0, "b":0, "c":21}
    {"_id":1, "a":0, "b":0, "c":54}
    .....
db.fun.aggregate([{$group: {_id: {a: "$a", b: "$b"},
                            c: {$max: "$c"}}},
                  {$group: {_id: "$_id.a", c: {$min: "$c"}}}])
```

$Project:

- remove keys
- Add keys
- Use some simple functions
  - $toUpper
  - $toLower
  - $add
  - $multiply.

Example:

```
db.Products.aggregate([
  { $Project:
    { _id:0 → exclude _id field
      'maker': { $toLower: "$manufacturer" },
      'details': { 'category': "$category",
                   'Price': { "$multiply": ["$price",
                 },
      'item': '$name' }}])
```

$Sort:

⇒ disk & memory based.
  → Default is in-memory based. (100 MB)

⇒ before or after the grouping stage.

Double unwind :

* → When there are two arrays in a document.
→ Can use $Push to reverse the impact of two $unwinds.

RDBMS	Mongo DB
WHERE	$match
HAVING	$match
SELECT	$project
ORDER BY	$sort
SUM()	$sum
JOIN	Mostly using $unwind.

limitations in Aggregation.

① 100 MB limit for Pipeline — Use allowDiskUse for this.
② 16 MB limit for document Size — Use Cursor = {...}
③ Sharded → Group by, Sort will cause it to work on the Primary Shard and not the replicas hence limiting Scalability when using Aggregation
⇒ Use Hadoop — Map/Reduce
Using mongo hadoop connector.

Week 6

Write Concern.

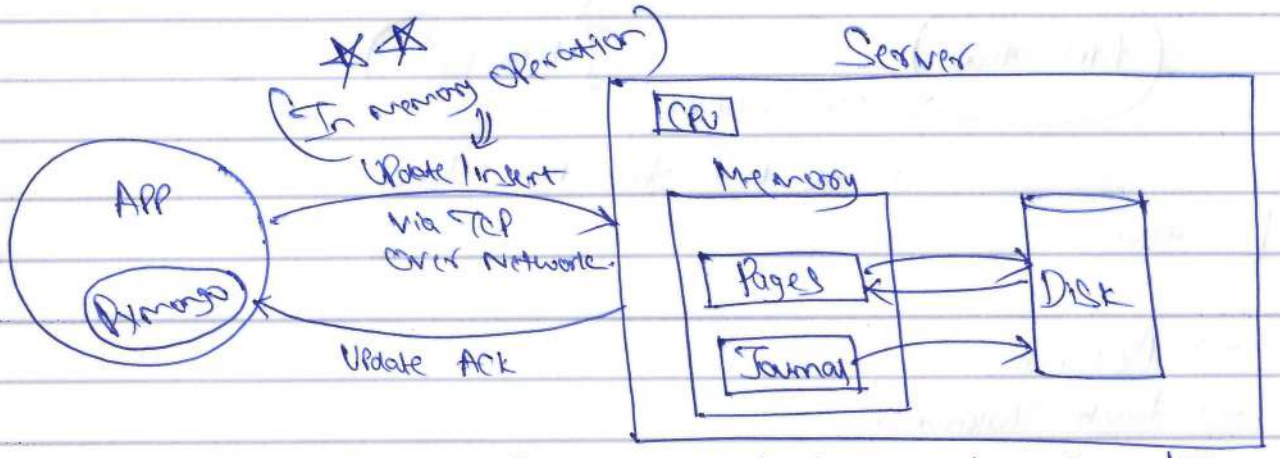

* Updates / Inserts first go and sits in the Pages / Journal area for a while before it goes to disk.
* Acknowledge ← write = 1, Journal = false ⟹ Write Concern.
* If there is a Server crash and if the Journal contents were not written to disk, then the data is lost.

w	j
1	false → fast, but small windows of vulnerability.
1	True → Slow and can be mentioned @ driver level.

Network Errors.

When w=1 & j=true

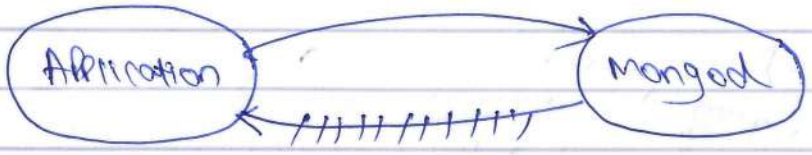

If no ACK then N/w Errors.

Replication

→ Availibility
→ Fault Tolerance.

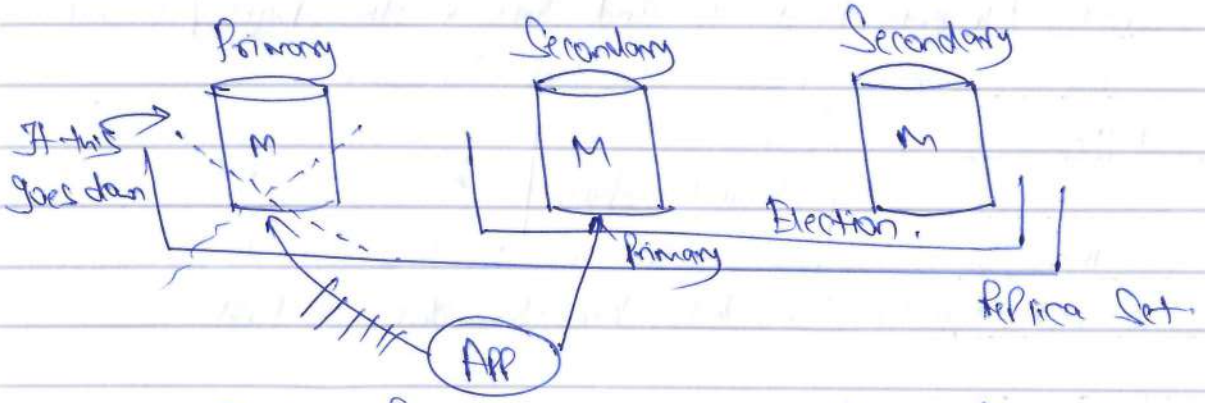

* If the Primary which went down comes back up, then this will join as the Secondary node
* Minimum number of nodes is 3. i.e to have the election Process.

* Replication is Asynchronous.

Types of Replica Set Nodes.

* Regular
* Arbiter – There for voting Purpose
* Delayed / Regular – Disaster Recovery node – Can't be
  ↳ Cannot be Primary (Priority = 0) in voting.
* Hidden – Used for Analytics. Priority = 0.

⇒ During Primary Down, no writes allowed.

Replication Internals.

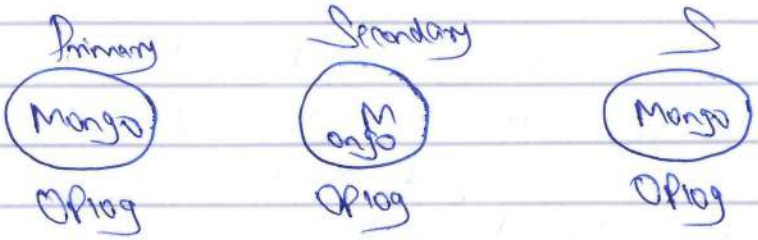

* Write on Primary will give information in Oplog and the Secondary will query this Oplog and apply the same Operation on itself.

Oplog: db.oplog.rs.find().pretty();

* Secondary will have the Optime and OptimeDate which tells us know the latest update Date / Time.
rs.getStatus(); ⇒ Check syncingTo also.

** If a node comes back up as a secondary after a period of being offline and the Oplog has looped / gone ahead on the primary, then
⇒ The entire data set will be copied from the primary.

Re-visiting write Concerns.

w=1 – wait for primary to Ack the write
w=2 – wait for primary & one other secondary
w=3 – ——— & all other secondary
wtime out – How long to wait?
→ Set on : ① Connection
② on a Collection
③ Replica Set → Safest to do.
w = majority → wait for majority of nodes to Ack.
** Journal waits for primary nodes only.

Read Preference

→ Default read Preference is Primary. To read from Secondary, Use rs.slaveOk();

→ Primary Preffered - Primary is Preffered, if not there then read from Secondary.

→ Eventually Consistent Read.

→ Secondary - Read from Secondary

→ Secondary Preffered - Secondary is Preffered, if not then Primary

→ Nearest - Based on min ping time.

Implication Of Replications:

* Seed lists - Node should be aware of other nodes
* Write Concern - w & j Parameter also wtimeout
* Read Preferences - Since multiple options are there
* Errors can happen - N/w errors, Exceptions, Syntax issues etc

SHARDING (Horizontal Scalability)

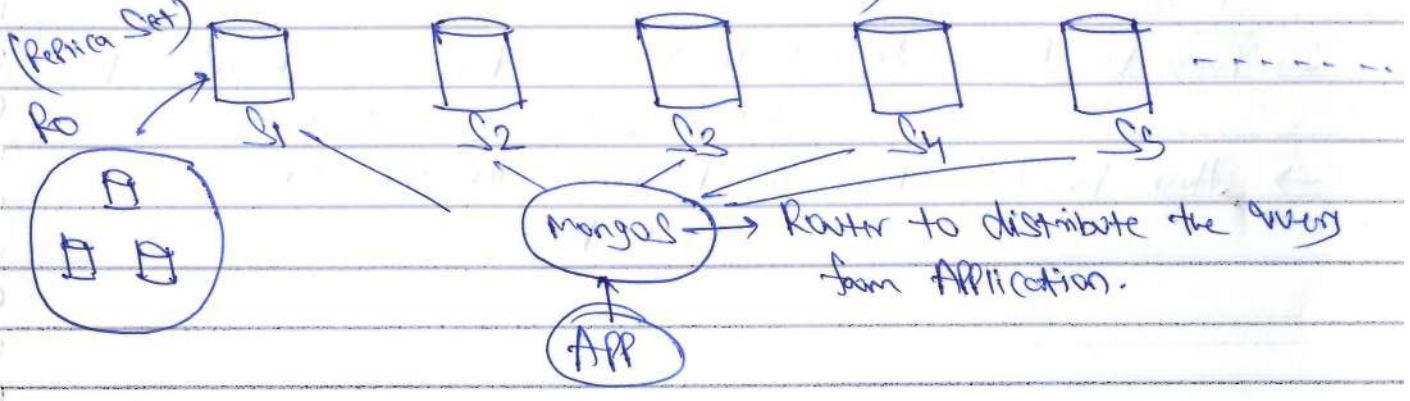

* Break up the collection into multiple logical halves and have a Shard key.

→ Range based Sharding
→ Shard key
Ex: Orders Table
* Breaked into chunks.

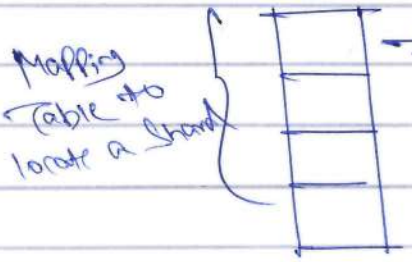

* Need to include Shard key on inserts.

* Application Queries the mongoS router which will find out a particular Shard where the request could be fulfilled.
* Multiple mongoS can be present.
* Each Shard is a replica Set again.

Sharding.
→ Range based – Ex: 0 to 100 in particular Shard S1 101 to 200 in S2 etc.
→ Hash based – Runs a hashing algorithm on the Shard key.

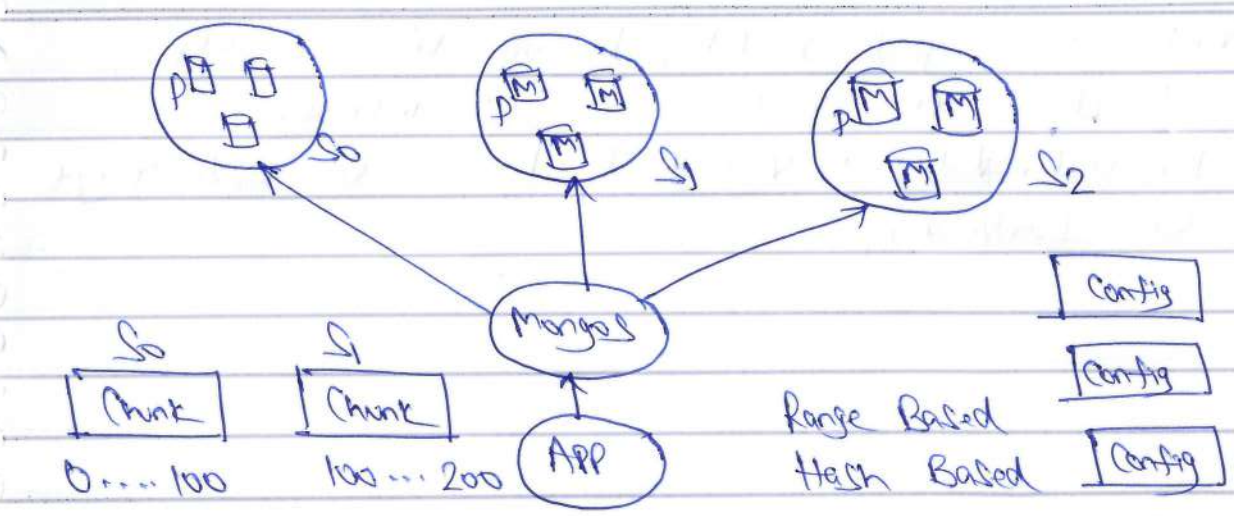

* Shard fail over → Mongos will try to re connect.

Choosing a Shard key.

* Sufficient Cardinality — i.e number of Possible values in a Column.
* Avoid Hot Spotting in Writes. i.e Problem of Writes always hitting the same Shard because of the Shard key. So choose a Shard key which is not just increasing in nature but have sufficient Cardinality

Ex: Orders

{ order_id ; ——,
order_date ; ——,
vendor ; —— }

① Vendor is a good Shard key because of Cardinality.
② Order_id is not great as its just increasing.
③ (Vendor, Order_date) is very good because of wide range of Cardinality.